

CISC 5597/6935 Distributed Systems

Lab 2: Consensus and Consistency— Raft Consensus Algorithm

Background:

Raft is a consensus algorithm used in distributed systems to ensure that data is replicated safely and consistently. It allows the members of a distributed system to agree on a sequence of values in the presence of failures.

Lab Assignment:

Based on RPC examples and the network model that you created in Lab-1, we are going to simulate a simplified distributed file system.

1. There are **THREE** “computing nodes” in your cluster. They are fully connected through RPC calls, meaning that they know each other's IPs and corresponding ports.
2. The cluster maintains a log file, **CISC6935**, that has replicas on each of the computing nodes.
3. It uses **RAFT** to maintain consistency on this simulated log file. The nodes have three states: Follower, Candidate, and Leader.

The cluster needs to simulate the following items.

1. Create a file on three nodes.
2. The client can submit MULTIPLE values, one at a time, to the system and the values will be stored in the log file distributedly.
 - a. For example, the client can call the RPC, SubmitValue, such as 1, on one of the nodes.
 - b. The system needs to write one line into the file, such as write 1.
 - c. When the system ensures consistency, it sends back a success message to the client.
3. Simulate the basic RAFT as discussed during lecture 8. There are 3 scenarios to be simulated.
 - a. Leader election phase and everything goes well, NO crashes or slow followers.
 - b. Simulate leader change in a perfect situation, which means NO left entries partially replicated (slide-8 page 54). It's basically the same as item a with another round of leader election, but picking the best candidate. This is a simplified selection due to the perfect situation.
 - i. The leader change process can start with the client, e.g., by typing a command
 - c. Simulate a scenario that the leader crashed (e.g., log file deleted) with an inconsistency log (e.g., missing only ONE log entry) on the other two followers. Check out slide-8 page 62.
 - i. Again, it can be started by the client's command.

- d. **Optional BONUS**: Simulate the crashed node rejoins the cluster. In this case, the leader needs to synchronize with the latest log files. The situation is similar to item c (slide-8 page 62), but with multiple rounds.
- e. Hint: You can use the sleep function to simulate network delay.

Grading Rubric: This lab allows a group of 2.

Your code should be elegant and well-documented (with comments).

- 1. The submission can not run successfully. (0 points)
- 2. Create a log file on three nodes. (10%)
- 3. The three computing nodes can communicate with each other through RPCs. (20%)
- 4. Leader election phase and everything goes well. (50%)
- 5. Item b. (10%)
- 6. Item c. (10%)
- 7. Bonus. (5%)