Lab 4
Naseem Machlovi

Option 3 (Maintenance): Monitoring a Kafka Cluster

Managing a cluster in real-time is an important task. Apache Kafka has several open-source monitoring tools, such as Kpow, Burrow, and Prometheus. In this lab, you are going to integrate a monitoring tool into your Kafka cluster and develop testing programs.

Requirements:

1. Deploy a cluster version of Kafka.
2. Ingrate a monitoring tool – choose any tool you want. Kpow and Burrow are preferred.
3. Develop producers and consumers that submit workloads to the system (you can find the code online).
4. Monitor your cluster with the tool that you installed, and you should see the workloads from the monitoring tool that you developed in Step 2.

1. Installation:

1. For Kafka installation, we use the following script on each node.

```
sudo apt-get update
sudo apt-get install default-jdk
wget https://downloads.apache.org/kafka/latest/kafka_2.13-2.8.0.tgz
tar -xzf kafka_2.13-2.8.0.tgz
cd kafka_2.13-2.8.0
mv kafka_2.13-2.8.0 kafka
```

2. Once installation is done make the following changes in zookeeper.properties on each node, to add our internal ip address in zookeeper.properties, to allow the server to listen each node

3. We have to create the following directories and give permission to the root. Make it on every node

```
<username>--> your username
<id> uniquebroker id [1,2,3]
sudo mkdir -p /var/lib/zookeeper
sudo chown <username>:<username> /var/lib/zookeeper
echo <id> | sudo tee /var/lib/zookeeper/myid
```

4. Now start the zookeeper on each node first and then start the kafka server

```
~/kafka/bin/zookeeper-server-start.sh -daemon ~/kafka/config/zookeeper.properties
```

~/kafka/bin/kafak-server-start.sh -daemon ~/kafka/config/server.properties
Now, we are ready to run the producer.py and consumer.py code on any client and can view the created topics on each node.

2. Kpow

We integrated Kpow as a monitoring tool for the Kafka cluster. Installation is followed by docker installation on the manager node pulling the latest image and running allowing to listen to port 3000.

Now open the following link in a tab after changing the firewall rules in GCP, allowing port TCP port 3000 to listen to external services.

http://<manager_external_ip>:3000

```
munee@worker3:~$ sudo docker run --pull=always -p 3000:3000 -m 2G factorhouse/kpow-ce:latest
latest: Pulling from factorhouse/kpow-ce
Digest: sha256:4f43289d52b86d328000063beeab0a5b55deece9784dd3771734bc7c9131c82b
Status: Image is up to date for factorhouse/kpow-ce:latest
```

Now you have to configure you community license, following the directions and adding the internal ips:9092 in bootstrap server and you can see the its monitoring

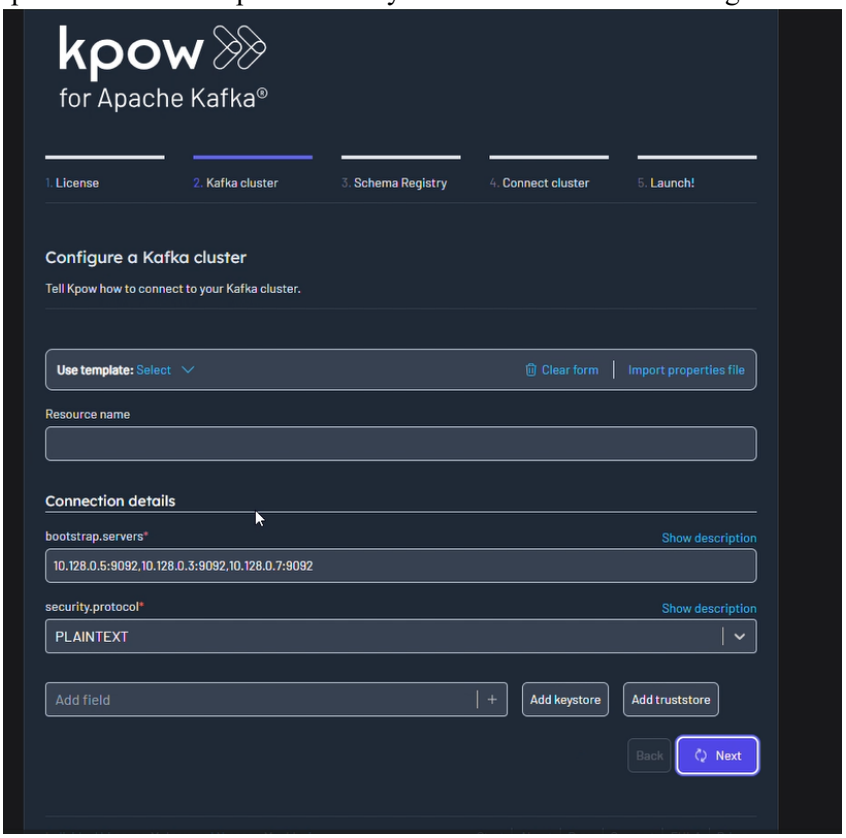

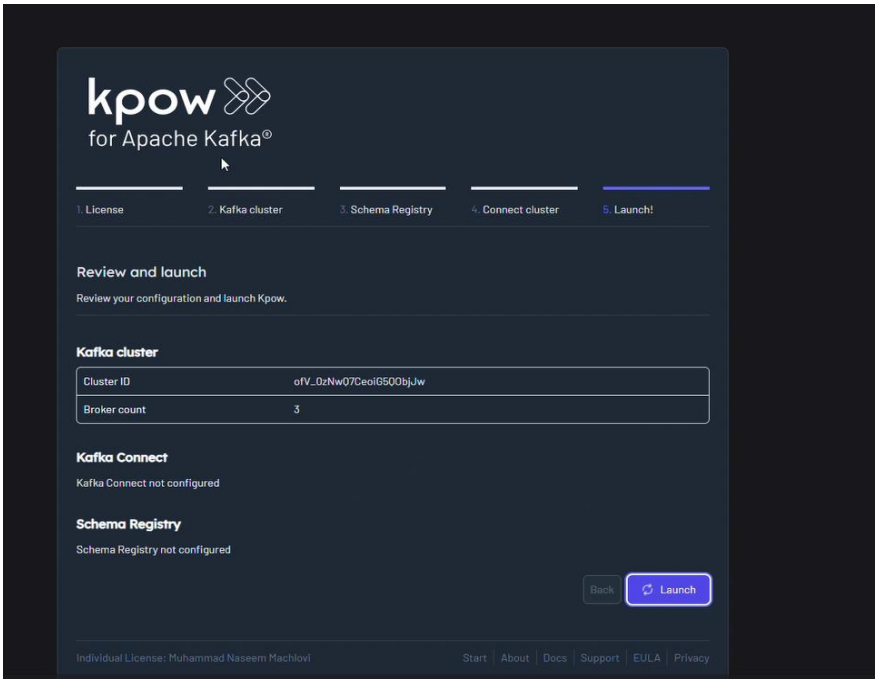
kpow
for Apache Kafka®
1. License
2. Kafka cluster
3. Schema Registry
4. Connect cluster
5. Launch!
Review and launch
Review your configuration and launch Kpow.
Kafka cluster
Cluster ID
ofV_0zNw07CeoiG5OObjJw
Broker count
3
Kafka Connect
Kafka Connect not configured
Schema Registry
Schema Registry not configured
Back
Launch
Individual License: Muhammad Naseem Machlovi
Start | About | Docs | Support | EULA | Privacy

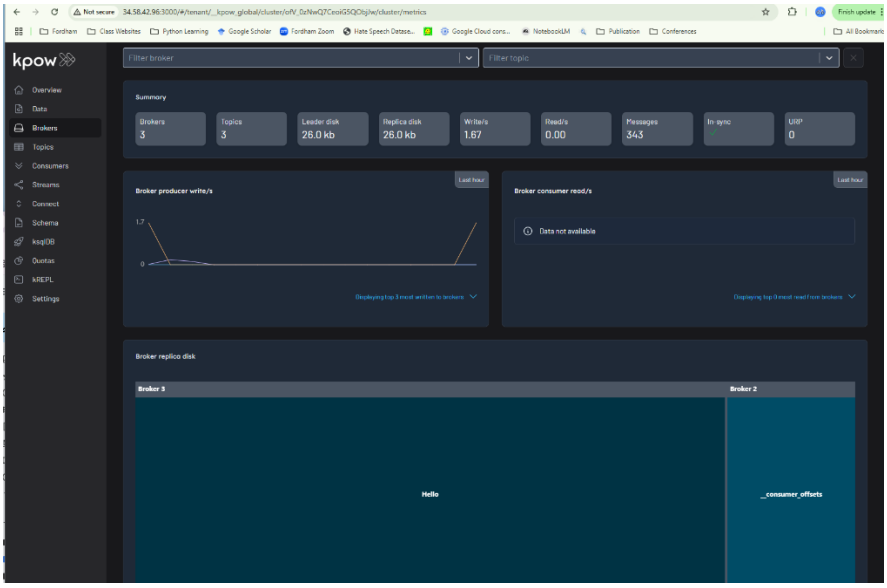
kpow
Filter broker
Filter topic
Summary
Brokers
3
Topics
3
Leader disk
26.0 kb
Replica disk
26.0 kb
Writes
1.67
Reads
0.00
Messages
343
In-sync
URP
0
Broker producer write/s
Broker consumer read/s
Data not available
Broker replica disk
Broker 3
Broker 2
Hello
_consumer_offsets

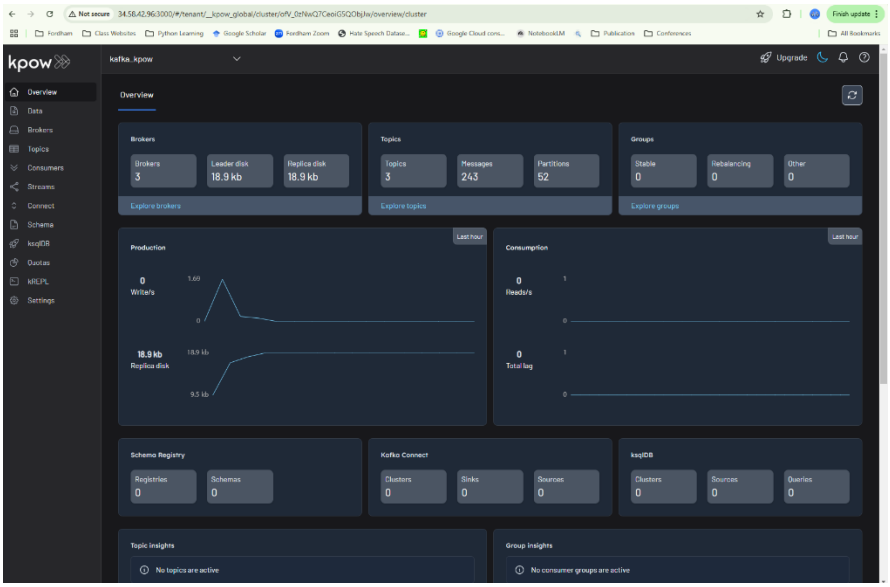
Not secure
34.56.42.96:3000/#/tenant/__kpow_global/cluster/ofV_GtNwQ7CeoiGSQObjJw/overview/cluster
Finish update
kpow
kafka_kpow
Upgrade
Overview
Data
Brokers
Topics
Consumers
Streams
Connect
Schema
ksqlDB
Quotas
kREPL
Settings
Overview
Brokers
Brokers 3
Leader disk 18.9 kb
Replica disk 18.9 kb
Explore brokers
Topics
Topics 3
Messages 243
Partitions 52
Explore topics
Groups
Stable 0
Rebalancing 0
Other 0
Explore groups
Production
Last hour
0
Write/s
1.69
0
18.9 kb
Replica disk
18.9 kb
9.5 kb
Consumption
Last hour
0
Reads/s
1
0
0
Total lag
1
0
Schema Registry
Registries 0
Schemas 0
Kafka Connect
Clusters 0
Sinks 0
Sources 0
ksqlDB
Clusters 0
Sources 0
Queries 0
Topic insights
No topics are active
Group insights
No consumer groups are active